



MIT 6.828: Operating System Engineering

Lab Report

Lab 1: C, Assembly, Tools, and Bootstrapping

<https://github.com/csnlp/MIT6828>

CSNLP

csnlp16@126.com

<https://csnlp.github.io/>

Contents

1	Environment Setup	2
2	Part 1: PC Bootstrap	3
2.1	Getting Started with x86 assembly	3
2.2	Simulating the x86	3
2.3	XX	3
2.4	The PC's Physical Address Space	3
2.5	The ROM BIOS	3
2.5.1	Exercise 2	3
3	Part 2: The Boot Loader	5
3.1	Loading The Kernel	6
3.1.1	What's ELF	6
3.1.2	Exercise 5	6
3.1.3	Exercise 6	6
4	Part 3: The Kernel	9
4.1	Using virtual memory to work around position dependence	9
4.1.1	Exercise 7	9
4.2	Formatted Printing to the Console	10
4.2.1	The Analysis of kern/printf.c, lib/printfmt.c, and kern/console.c	10
4.2.2	Analysis of kern/console.c	10
4.2.3	Analysis of lib/printfmt.c	10
4.2.4	Exercise 8	10
4.3	The Questions Following Exercise 8	14
4.4	The Stack	17
4.4.1	The ESP and EBP Register	17
4.4.2	Exercise 9	17
4.4.3	Exercise 10	20
5	Reference	22

1 Environment Setup

Our environment: 1. Host: Mac Air 2. Virtualbox: Ubuntu 14.04

Install The QEMU MIT6.828 uses QEMU as simulator. To install QEMU, we use the command like: `sudo apt-get install qemu`

<https://github.com/csnlp/MIT6828>

2 Part 1: PC Bootstrap

Our environment: 1. Host: Mac Air 2. Virtualbox: Ubuntu 14.04

Install The QEMU MIT6.828 uses QEMU as simulator. To install QEMU, we use the command like: `sudo apt-get install qemu`

2.1 Getting Started with x86 assembly

- Instructions classes:
 - Data movement: MOV, PUSH, POP, ...
 - Arithmetic: TEST, SHL, ADD, ...
 - I/O: IN, OUT, ...
 - Control: JMP, JZ, JNZ, CALL, RET
 - String: REP, MOVSB, ...
 - System: IRET, INT, ...
- Intel architecture manual Volume 2
 - Intel syntax: `op dst, src`
 - AT&T (gcc/gas) syntax: `op src, dst`

Figure 1. This is a caption

2.2 Simulating the x86

2.3 XX

2.4 The PC's Physical Address Space

I believe PC physical address space is considerably important.

2.5 The ROM BIOS

2.5.1 Exercise 2

Use GDB's `si` (Step Instruction) command to trace into the ROM BIOS for a few more instructions, and try to guess what it might be doing. You might want to look at Phil Storrs I/O Ports Description, as well as other materials on the 6.828 reference materials page. No need to figure out all the details - just the general idea of what the BIOS is doing first.

Here is the experiment steps:

1. Open a terminal and type: `make qemu-gdb`.

2. Open another terminal and type: `make gdb`

We can see the output as follows:

```
The target architecture is assumed to be i8086
[f000:fff0] 0xffff0: ljmp  $0xf000,$0xe05b
0x0000fff0 in ?? ()
+ symbol-file obj/kern/kernel
(gdb) si
[f000:e05b] 0xfe05b: cml  $0x0,%cs:0x6574
0x0000e05b in ?? ()
(gdb) █
```

Figure 2: This is a caption

<https://github.com/csnlp/MIT6828>

3 Part 2: The Boot Loader

The boot loader must do

1. Switchs from the real-mode to 32-bit protected mode.
2. Reads the kernel from the hard disk by directly accessing the IDE disk device registers via the x86's special I/O instructions.

The first assembly instruction is

```
1 [f000: fff0 ] 0xffff0 : ljmp $0xf000 , $0xe05b
```

Listing 1: The first assembly instruction of boot program

The first instruction is to jump from the **real mode** to **protective mode**. With [si](#) instruction, we can see the following instructions:

```
1 [f000:e05b] 0xfe05b: cmpl    $0x0,%cs:0x6574
2 [f000:e062] 0xfe062: jne     0xfd2b6
3 [f000:e066] 0xfe066: xor     %ax,%ax
4 [f000:e068] 0xfe068: mov     %ax,%ss
5 [f000:e06a] 0xfe06a: mov     $0x7000,%esp
6 [f000:e070] 0xfe070: mov     $0xf3c24,%edx
7 [f000:e076] 0xfe076: jmp     0xfd124
8 [f000:d124] 0xfd124: mov     %eax,%ecx
9 [f000:d127] 0xfd127: cli
10 [f000:d128] 0xfd128: cld
11 [f000:d129] 0xfd129: mov     $0x8f,%eax
12 [f000:d12f] 0xfd12f: out     %al,$0x70
13 [f000:d131] 0xfd131: in      $0x71,%al
14 [f000:d133] 0xfd133: in      $0x92,%al
15 [f000:d135] 0xfd135: or      $0x2,%al
```

Listing 2: The first assembly instruction of boot program

The question asked in

1. At what point does the processor start executing 32-bit code? What exactly causes the switch from 16- to 32-bit mode?

Answer: we can see in [boot/boot.S](#) that

```
1 ljmp     $PROT_MODE_CSEG, $protcseg
```

Listing 3: The switch from 16 to 32 bit mode

2. What is the last instruction of the boot loader executed, and what is the first instruction of the kernel it just loaded?

Answer: it the command in [boot/main.c](#)

```
1 ((void (*)(void)) (ELFHDR->e_entry))();
```

Listing 4: Last instruction of boot loader

Additional Information:

3. Where is the first instruction of the kernel?
4. How does the boot loader decide how many sectors it must read in order to fetch the entire kernel from disk? Where does it find this information?

3.1 Loading The Kernel

3.1.1 What's ELF

3.1.2 Exercise 5

[Exercise 5. Trace through the first few instructions of the boot loader again and identify the first instruction that would "break" or otherwise do the wrong thing if you were to get the boot loader's link address wrong. Then change the link address in `boot/Makefrag` to something wrong, run `make clean`, recompile the lab with `make`, and trace into the boot loader again to see what happens. Don't forget to change the link address back and `make clean` again afterward!]

We can see from `boot/Makefrag` that

```
1 $(OBJDIR)/boot/boot: $(BOOT_OBJS)
2     @echo + ld boot/boot
3     $(V)$(LD) $(LDFLAGS) -N -e start -Ttext 0x7C00 -o $@.out $^
4     $(V)$(OBJDUMP) -S $@.out >$@.asm
5     $(V)$(OBJCOPY) -S -O binary -j .text $@.out $@
6     $(V)perl boot/sign.pl $(OBJDIR)/boot/boot
```

Listing 5: makefrag

We do the following operations

1. Open a terminal and type `make qemu-gdb`.
2. Open a second terminal and type `make gdb`
3. Set a breakpoint and continue execution to this breakpoint: type in the second terminal that `b *0x7c00` and `c`
4. Repeat the command `si` in the second terminal.

Close these two terminal.

Then, run `make clean`. After that change the address in the `boot/Makefrag` to `0x7d00` and then `make`.

1. Open a terminal and type `make qemu-gdb`
2. Open a second terminal and type `make gdb`
3. Set a breakpoint and continue execution to this breakpoint: type in the second terminal that `b *0x7c00` and `c`
4. Repeat the command `si` in the second terminal.

We can see that:

3.1.3 Exercise 6

[We can examine memory using GDB's `x` command. The GDB manual has full details, but for now, it is enough to know that the command `x/Nx ADDR` prints `N` words of memory at `ADDR`. (Note that both 'x's in the command are lowercase.) Warning: The size of a word is not a universal standard. In GNU assembly, a word is two bytes (the 'w' in `xorw`, which stands for word, means 2 bytes).

```

(gdb) si
[ 0:7c1a] => 0x7c1a:  mov    $0xdf,%al
0x00007c1a in ?? ()
(gdb) si
[ 0:7c1c] => 0x7c1c:  out    %al,$0x60
0x00007c1c in ?? ()
(gdb) si
[ 0:7c1e] => 0x7c1e:  lgdtw  0x7c64
0x00007c1e in ?? ()
(gdb) si
[ 0:7c23] => 0x7c23:  mov    %cr0,%eax
0x00007c23 in ?? ()
(gdb) si
[ 0:7c26] => 0x7c26:  or     $0x1,%eax
0x00007c26 in ?? ()
(gdb) si
[ 0:7c2a] => 0x7c2a:  mov    %eax,%cr0
0x00007c2a in ?? ()
(gdb) si
[ 0:7c2d] => 0x7c2d:  ljmp   $0x8,$0x7c32
0x00007c2d in ?? ()
(gdb) si
The target architecture is assumed to be i386
=> 0x7c32:      mov    $0x10,%ax
0x00007c32 in ?? ()
(gdb) si
=> 0x7c36:      mov    %eax,%ds
0x00007c36 in ?? ()
(gdb) 

```

Figure 3: The correct version with the link address as 0x7c00


```

(gdb) si
[ 0:7c1a] => 0x7c1a: mov    $0xdf,%al
0x00007c1a in ?? ()
(gdb) si
[ 0:7c1c] => 0x7c1c: out    %al,$0x60
0x00007c1c in ?? ()
(gdb) si
[ 0:7c1e] => 0x7c1e: lgdtw  0x7d64
0x00007c1e in ?? ()
(gdb) si
[ 0:7c23] => 0x7c23: mov    %cr0,%eax
0x00007c23 in ?? ()
(gdb) si
[ 0:7c26] => 0x7c26: or     $0x1,%eax
0x00007c26 in ?? ()
(gdb) si
[ 0:7c2a] => 0x7c2a: mov    %eax,%cr0
0x00007c2a in ?? ()
(gdb) si
[ 0:7c2d] => 0x7c2d: ljmp   $0x8,$0x7d32
0x00007c2d in ?? ()
(gdb) si
[f000:e05b] 0xfe05b: cmpl   $0x0,%cs:0x6574
0x0000e05b in ?? ()
(gdb) si
[f000:e062] 0xfe062: jne    0xfd2b6
0x0000e062 in ?? ()
(gdb)

```

Figure 4: The wrong version with the link address as 0x7d00

Reset the machine (exit QEMU/GDB and start them again). Examine the 8 words of memory at 0x00100000 at the point the BIOS enters the boot loader, and then again at the point the boot loader enters the kernel. Why are they different? What is there at the second breakpoint? (You do not really need to use QEMU to answer this question. Just think.)]

We can clearly observe that after the kernel was loaded (**b *0x1000c** and **c**), the memory **0x100000** was loaded with kernel code.

```

(gdb) b *0x7c00
Breakpoint 1 at 0x7c00
(gdb) c
Continuing.
[ 0:7c00] => 0x7c00: cli

Breakpoint 1, 0x00007c00 in ?? ()
(gdb) x/8x 100000
0x186a0: 0x00000000 0x00000000 0x00000000 0x00000000
0x186b0: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) x/8x 0x100000
0x100000: 0x00000000 0x00000000 0x00000000 0x00000000
0x100010: 0x00000000 0x00000000 0x00000000 0x00000000
(gdb) b *0x10000c
Breakpoint 2 at 0x10000c
(gdb) c
Continuing.
The target architecture is assumed to be i386
=> 0x10000c: movw $0x1234,0x472

Breakpoint 2, 0x0010000c in ?? ()
(gdb) x/8x 0x100000
0x100000: 0x1badb002 0x00000000 0xe4524ffe 0x7205c766
0x100010: 0x34000004 0x0000b812 0x220f0011 0xc0200fd8
(gdb)

```

Figure 5: Exercise 6

4 Part 3: The Kernel

4.1 Using virtual memory to work around position dependence

4.1.1 Exercise 7

[Exercise 7. Use QEMU and GDB to trace into the JOS kernel and stop at the `movl %eax, %cr0`. Examine memory at `0x00100000` and at `0xf0100000`. Now, single step over that instruction using the `stepi` GDB command. Again, examine memory at `0x00100000` and at `0xf0100000`. Make sure you understand what just happened. What is the first instruction after the new mapping is established that would fail to work properly if the mapping weren't in place? Comment out the `movl %eax, %cr0` in `kern/entry.S`, trace into it, and see if you were right.]

Okay, first we have observe that `0x10000C` is the entry address of kernel. Let's firstly check the content in memory `0x00100000` and at `0xf0100000` before the instruction: `movl %eax, %cr0`.

```

(gdb) x/4xb 0xf0100000
0xf0100000 <_start+4026531828>: 0x00 0x00 0x00 0x00
(gdb) x/4xb 0x00100000
0x100000: 0x02 0xb0 0xad 0x1b

```

Figure 6: Exercise 7: After the instruction

After this instruction: `movl %eax, %cr0`, we can see that

```
(gdb) x/4xb 0xf0100000
0xf0100000 <_start+4026531828>: 0x02    0xb0    0xad    0x1b
(gdb) x/4xb 0x00100000
0x100000:    0x02    0xb0    0xad    0x1b
(gdb)
```

Figure 7: Exercise 7: After the instruction

4.2 Formatted Printing to the Console

4.2.1 The Analysis of kern/printf.c, lib/printfmt.c, and kern/console.c

We have the following observations:

1. [kern/printf.c](#) has the following two functions (it has have a static method `putch()` for its self.):
 - `cprintf()`: which use the `vprintfmt()` in file [lib/printfmt.c](#)
 - `vcprintf()`: use the `vcprintf()` in the [lib/printfmt.c](#)
 - `putch()`: use the `cputchar()` in the [kern/console.c](#).
2. `dd`

So, let's start from [kern/console.c](#).

4.2.2 Analysis of kern/console.c

```
1 void
2 cputchar(int c)
3 {
4     cons_putc(c);
5 }
```

Listing 6: The cputchar function in kern/console.c

Then, let's see `cons_putc` function

```
1 // output a character to the console
2 static void
3 cons_putc(int c)
4 {
5     serial_putc(c);
6     lpt_putc(c);
7     cga_putc(c);
8 }
```

Listing 7: The cons_putc function in kern/console.c

4.2.3 Analysis of lib/printfmt.c

4.2.4 Exercise 8

Exercise 8. We have omitted a small fragment of code - the code necessary to print octal numbers using patterns of the form `"%o"`. Find and fill in this code fragment.

```

1 // Main function to format and print a string.
2 void printfmt(void (*putch)(int, void*), void *putdat, const char *fmt, ...);
3
4 void
5 vprintfmt(void (*putch)(int, void*), void *putdat, const char *fmt, va_list ap)
6 {
7     register const char *p;
8     register int ch, err;
9     unsigned long long num;
10    int base, lflag, width, precision, altflag;
11    char padc;
12
13    while (1) {
14        //print all the characters before %
15        while ((ch = *(unsigned char *) fmt++) != '%') {
16            if (ch == '\\0')
17                return;
18            putch(ch, putdat);
19        }
20
21        // Process a %-escape sequence
22        padc = ' ';
23        width = -1;
24        precision = -1;
25        lflag = 0;
26        altflag = 0;
27    reswitch:
28        switch (ch = *(unsigned char *) fmt++) {
29
30            // flag to pad on the right
31            case '-':
32                padc = '-';
33                goto reswitch;
34
35            // flag to pad with 0's instead of spaces
36            case '0':
37                padc = '0';
38                goto reswitch;
39
40            // width field
41            case '1':
42            case '2':
43            case '3':
44            case '4':
45            case '5':
46            case '6':
47            case '7':
48            case '8':
49            case '9':
50                for (precision = 0; ; ++fmt) {
51                    precision = precision * 10 + ch - '0';
52                    ch = *fmt;
53                    if (ch < '0' || ch > '9')
54                        break;
55                }
56                goto process_precision;
57
58            case '*':

```

```

59     precision = va_arg(ap, int);
60     goto process_precision;
61
62 case '.':
63     if (width < 0)
64         width = 0;
65     goto reswitch;
66
67 case '#':
68     altflag = 1;
69     goto reswitch;
70
71 process_precision:
72     if (width < 0)
73         width = precision, precision = -1;
74     goto reswitch;
75
76 // long flag (doubled for long long)
77 case 'l':
78     lflag++;
79     goto reswitch;
80
81 // character
82 case 'c':
83     putchar(va_arg(ap, int), putdat);
84     break;
85
86 // error message
87 case 'e':
88     err = va_arg(ap, int);
89     if (err < 0)
90         err = -err;
91     if (err >= MAXERROR || (p = error_string[err]) == NULL)
92         printfmt(putch, putdat, "error %d", err);
93     else
94         printfmt(putch, putdat, "%s", p);
95     break;
96
97 // string
98 case 's':
99     if ((p = va_arg(ap, char *)) == NULL)
100         p = "(null)";
101     if (width > 0 && padc != '-')
102         for (width -= strlen(p, precision); width > 0; width--)
103             putchar(padc, putdat);
104     for (; (ch = *p++) != '\0' && (precision < 0 || --precision >= 0); width--)
105         if (altflag && (ch < ' ' || ch > '~'))
106             putchar('?', putdat);
107         else
108             putchar(ch, putdat);
109     for (; width > 0; width--)
110         putchar(' ', putdat);
111     break;
112
113 // (signed) decimal
114 case 'd':
115     num = getint(&ap, lflag);
116     if ((long long) num < 0) {

```

```

117     putchar('-', putdat);
118     num = -(long long) num;
119 }
120 base = 10;
121 goto number;
122
123 // unsigned decimal
124 case 'u':
125     num = getuint(&ap, lflag);
126     base = 10;
127     goto number;
128
129 // (unsigned) octal
130 case 'o':
131     // Replace this with your code.
132     putchar('X', putdat);
133     putchar('X', putdat);
134     putchar('X', putdat);
135     break;
136
137 // pointer
138 case 'p':
139     putchar('0', putdat);
140     putchar('x', putdat);
141     num = (unsigned long long)
142         (uintptr_t) va_arg(ap, void *);
143     base = 16;
144     goto number;
145
146 // (unsigned) hexadecimal
147 case 'x':
148     num = getuint(&ap, lflag);
149     base = 16;
150 number:
151     printnum(putch, putdat, num, base, width, padc);
152     break;
153
154 // escaped '%' character
155 case '%':
156     putchar(ch, putdat);
157     break;
158
159 // unrecognized escape sequence - just print it literally
160 default:
161     putchar('%', putdat);
162     for (fmt--; fmt[-1] != '%'; fmt--)
163         /* do nothing */;
164     break;
165 }
166 }
167 }

```

Listing 8: The vprintfmt() function in lib/printfmt.c

1. Print all the character before %.
2. Then it process the special character just before % character:

```

while ((ch = *(unsigned char *) fmt++) != '%') {
    if (ch == '\\0')
        return;
    putchar(ch, putdat);
}

```

Figure 8: Exercise 7: After the instruction

- long: **l**
- character: **c**
- error: **e**
- decimal: **d**
- ...

Actually, **case 'o'** is where we should fill in our code:

```

1 // (unsigned) octal
2 case 'o':
3     // Replace this with your code.
4     //putch('X', putdat);
5     //putch('X', putdat);
6     //putch('X', putdat);
7     //break;
8     num = getuint(&ap, lflag);
9     base = 8;
10    goto number;

```

Listing 9: Case octal

4.3 The Questions Following Exercise 8

1. *[Explain the interface between printf.c and console.c. Specifically, what function does console.c export? How is this function used by printf.c?]*

Answer: **console.c** export the **cputchar()** function.

2. *[Explain the following from console.c:]*

```

1 // What is the purpose of this?
2 if (crt_pos >= CRT_SIZE) {
3     int i;
4
5     memmove(crt_buf, crt_buf + CRT_COLS, (CRT_SIZE - CRT_COLS) *
6     ↪ sizeof(uint16_t));
7     for (i = CRT_SIZE - CRT_COLS; i < CRT_SIZE; i++)
8         crt_buf[i] = 0x0700 | ' ';
9     crt_pos -= CRT_COLS;
10 }

```

Listing 10: Subpart of console.c

The `memmove`¹ command is used to put the memory at `crt_buf + CRT_COLS` to `crt_buf`. This reason is to solve the problem that the current `crt_pos` is already bigger than current `CRT_SIZE`. After that

3. For the following questions you might wish to consult the notes for Lecture 2. These notes cover GCC's calling convention on the x86. Trace the execution of the following code step-by-step:

```
1  int x = 1, y = 3, z = 4;
2  cprintf("x %d, y %x, z %d\n", x, y, z);
3
```

Listing 11: question 3 code

- In the call to `printf()`, to what does `fmt` point? To what does `ap` point?
- List (in order of execution) each call to `cons_putc`, `va_arg`, and `vcprintf`. For `cons_putc`, list its argument as well. For `va_arg`, list what `ap` points to before and after the call. For `vcprintf` list the values of its two arguments.

Answer: Let's first revisit the `vcprintf()` in `kern/printf.c`:

```
1  static void
2  putch(int ch, int *cnt)
3  {
4      cputchar(ch);
5      *cnt++;
6  }
7
8  int
9  vcprintf(const char *fmt, va_list ap)
10 {
11     int cnt = 0;
12
13     vprintfmt((void*)putch, &cnt, fmt, ap);
14     return cnt;
15 }
16
17 int
18 cprintf(const char *fmt, ...)
19 {
20     va_list ap;
21     int cnt;
22
23     va_start(ap, fmt);
24     cnt = vcprintf(fmt, ap);
25     va_end(ap);
26
27     return cnt;
28 }
29
```

Listing 12: question 3 code

¹`memmove(void destination, const void source, size_t num)`

We can see that `fmt` is the pointer to this string. `va_list ap` is the list of the input parameter in the string `fmt`, i.e., `{x, y, z}`.

We can see that `cprintf()` calls `vcprintf()`, then the `vcprintf()` call the `vprintfmt()`. During the process of `vprintfmt()`, it jumps to the `case d`.

```

1      case 'd':
2          num = getint(&ap, lflag);
3          if ((long long) num < 0) {
4              putchar('-', putdat);
5              num = -(long long) num;
6          }
7          base = 10;
8          goto number;

```

We can see that it first output the minus symbol '-' if number is negative. Then `goto number`. Actually, `goto number` is for

- `case 'd'`
- `case 'u'`
- `case 'p'`

4. Run the following code

```

1      unsigned int i = 0x00646c72;
2      cprintf("H%x Wo%s", 57616, &i);

```

Listing 13: Question 4 Code

What is the output? Explain how this output is arrived at in the step-by-step manner of the previous exercise. Here's an ASCII table that maps bytes to characters. The output depends on that fact that the x86 is little-endian. If the x86 were instead big-endian what would you set `i` to in order to yield the same output? Would you need to change 57616 to a different value?

Answer: It print `He110 World`

5. In the following code, what is going to be printed after `'y='`? (note: the answer is not a specific value.) Why does this happen? `cprintf("x=%d y=%d", 3);`

Answer: We append the line in the `kern/monitor.c` file and find the result as: In our case

```

Welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
x=3, y=-267380388
K>

```

Figure 9: Question 5

4.4 The Stack

4.4.1 The ESP and EBP Register

1. ESP

- The x86 stack pointer (esp register) points to the lowest location on the stack that is currently in use. Everything below that location in the region reserved for the stack is free.
- **Pushing** a value onto the stack involves decreasing the stack pointer and then writing the value to the place the stack pointer points to.
- **Popping** a value from the stack involves reading the value the stack pointer points to and then increasing the stack pointer.

2. EBP

- d

4.4.2 Exercise 9

[Exercise 9. Determine where the kernel initializes its stack, and exactly where in memory its stack is located. How does the kernel reserve space for its stack? And at which "end" of this reserved area is the stack pointer initialized to point to?]

Answer: we need to carefully read the [kern/entry.S](#) file.

```
1 /* See COPYRIGHT for copyright information. */
2
3 #include <inc/mmu.h>
4 #include <inc/memlayout.h>
5
6 # Shift Right Logical
7 #define SRL(val, shamt) (((val) >> (shamt)) & ~(-1 << (32 - (shamt))))
8
9
10 #####
11 # The kernel (this code) is linked at address ~(KERNBASE + 1 Meg),
12 # but the bootloader loads it at address ~1 Meg.
13 #
14 # RELOC(x) maps a symbol x from its link address to its actual
15 # location in physical memory (its load address).
16 #####
17
18 #define RELOC(x) ((x) - KERNBASE)
19
20 #define MULTIBOOT_HEADER_MAGIC (0x1BADB002)
21 #define MULTIBOOT_HEADER_FLAGS (0)
22 #define CHECKSUM (-(MULTIBOOT_HEADER_MAGIC + MULTIBOOT_HEADER_FLAGS))
23
24 #####
25 # entry point
26 #####
27
28 .text
29
30 # The Multiboot header
31 .align 4
```

```

32 .long MULTIBOOT_HEADER_MAGIC
33 .long MULTIBOOT_HEADER_FLAGS
34 .long CHECKSUM
35
36 # '_start' specifies the ELF entry point. Since we haven't set up
37 # virtual memory when the bootloader enters this code, we need the
38 # bootloader to jump to the *physical* address of the entry point.
39 .globl _start
40 _start = RELOC(entry)
41
42 .globl entry
43 entry:
44     movw $0x1234,0x472    # warm boot
45
46     # We haven't set up virtual memory yet, so we're running from
47     # the physical address the boot loader loaded the kernel at: 1MB
48     # (plus a few bytes). However, the C code is linked to run at
49     # KERNBASE+1MB. Hence, we set up a trivial page directory that
50     # translates virtual addresses [KERNBASE, KERNBASE+4MB) to
51     # physical addresses [0, 4MB). This 4MB region will be
52     # sufficient until we set up our real page table in mem_init
53     # in lab 2.
54
55     # Load the physical address of entry_pgdir into cr3. entry_pgdir
56     # is defined in entrypgdir.c.
57     movl $(RELOC(entry_pgdir)), %eax
58     movl %eax, %cr3
59     # Turn on paging.
60     movl %cr0, %eax
61     orl $(CRO_PE|CRO_PG|CRO_WP), %eax
62     movl %eax, %cr0
63
64     # Now paging is enabled, but we're still running at a low EIP
65     # (why is this okay?). Jump up above KERNBASE before entering
66     # C code.
67     mov $relocated, %eax
68     jmp *%eax
69 relocated:
70
71     # Clear the frame pointer register (EBP)
72     # so that once we get into debugging C code,
73     # stack backtraces will be terminated properly.
74     movl $0x0,%ebp    # nuke frame pointer
75
76     # Set the stack pointer
77     movl $(bootstacktop),%esp
78
79     # now to C code
80     call i386_init
81
82     # Should never get here, but in case we do, just spin.
83 spin: jmp spin
84
85
86 .data
87 #####
88 # boot stack
89 #####

```

```

90 .p2align PGSHIFT # force page alignment
91 .globl bootstack
92 bootstack:
93 .space KSTKSIZE
94 .globl bootstacktop
95 bootstacktop:

```

Listing 14: kern/entry.S

We can see that

```

1  movl $0x0,%ebp # nuke frame pointer
2  # Set the stack pointer
3  movl $(bootstacktop),%esp

```

is used to initialize its stack.

1. Determine where the kernel initializes its stack? **Answer:** in the [kern/entry.S](#) file. The `movl $0x0,%ebp` and `movl $(bootstacktop),%esp` command.

2. Where in memory its stack is located?

To better understand this problem, we should separate it into the following steps: we give the instruction in [kern/entry.S](#) again.

```

1      # Load the physical address of entry_pgdire into cr3.  entry_pgdire
2      # is defined in entrypgdire.c.
3      movl $(RELOC(entry_pgdire)), %eax
4      movl %eax, %cr3
5      # Turn on paging.
6      movl %cr0, %eax
7      orl $(CR0_PE|CR0_PG|CR0_WP), %eax
8      movl %eax, %cr0
9
10     # Now paging is enabled, but we're still running at a low EIP
11     # (why is this okay?). Jump up above KERNBASE before entering
12     # C code.
13     mov $relocated, %eax
14     jmp *%eax
15 relocated:
16
17     # Clear the frame pointer register (EBP)
18     # so that once we get into debugging C code,
19     # stack backtraces will be terminated properly.
20     movl $0x0,%ebp # nuke frame pointer
21
22     # Set the stack pointer
23     movl $(bootstacktop),%esp
24
25     # now to C code
26     call i386_init
27

```

- Load the physical address of table `entry_pgdire` to `cr3`.
- Set the PE, PG, WP of `cr0` as 1.
- Jump up above `KERNBASE` before entering before entering C code. The `$relocated` value is `0xf010002`

- Set the **ebp** as **0x00**.
- Set the **esp** as **0xf0110000**.

The stack size is 32KB. So, the stack locates **0xf0108000-0xf0111000**

4.4.3 Exercise 10

[Exercise 10. To become familiar with the C calling conventions on the x86, find the address of the test_backtrace function in obj/kern/kernel.asm, set a breakpoint there, and examine what happens each time it gets called after the kernel starts. How many 32-bit words does each recursive nesting level of test_backtrace push on the stack, and what are those words?]

Note that, for this exercise to work properly, you should be using the patched version of QEMU available on the tools page or on Athena. Otherwise, you'll have to manually translate all breakpoint and memory addresses to linear addresses.

Answer: Let's firstly check the C code of **test_backtrace** in **kern/init.c**:

```
1 // Test the stack backtrace function (lab 1 only)
2 void
3 test_backtrace(int x)
4 {
5     cprintf("entering test_backtrace %d\n", x);
6     if (x > 0)
7         test_backtrace(x-1);
8     else
9         mon_backtrace(0, 0, 0);
10    cprintf("leaving test_backtrace %d\n", x);
11 }
```

From this code, we can figure out why these code when we **make qemu**.

```
cui@cui-VirtualBox:~/mit6828/lab$ make qemu
sed "s/localhost:1234/localhost:26000/" < .gdbinit.tmpl > .gdbinit
qemu-system-i386 -drive file=obj/kern/kernel.img,index=0,media=disk,format=raw -
serial mon:stdio -gdb tcp::26000 -D qemu.log
6828 decimal is 15254 octal!
entering test_backtrace 5
entering test_backtrace 4
entering test_backtrace 3
entering test_backtrace 2
entering test_backtrace 1
entering test_backtrace 0
leaving test_backtrace 0
leaving test_backtrace 1
leaving test_backtrace 2
leaving test_backtrace 3
leaving test_backtrace 4
leaving test_backtrace 5
welcome to the JOS kernel monitor!
Type 'help' for a list of commands.
K>
```

Figure 10: Exercise 6

Because there is **test_backtrace(5)** in the **kern/init.c**.

Then we check the **test_backtrace** in **obj/kern/kernel.asm**

```

1 // Test the stack backtrace function (lab 1 only)
2 void
3 test_backtrace(int x)
4 {
5 f0100040:      55                push    %ebp
6 f0100041:      89 e5            mov     %esp,%ebp
7 f0100043:      53                push    %ebx
8 f0100044:      83 ec 14         sub     $0x14,%esp
9 f0100047:      8b 5d 08         mov     0x8(%ebp),%ebx
10      cprintf("entering test_backtrace %d\n", x);
11 f010004a:      89 5c 24 04      mov     %ebx,0x4(%esp)
12 f010004e:      c7 04 24 e0 18 10 f0 movl    $0xf01018e0,(%esp)
13 f0100055:      e8 d7 08 00 00   call    f0100931 <cstdio>
14      if (x > 0)
15 f010005a:      85 db            test    %ebx,%ebx
16 f010005c:      7e 0d            jle     f010006b <test_backtrace+0x2b>
17      test_backtrace(x-1);
18 f010005e:      8d 43 ff         lea     -0x1(%ebx),%eax
19 f0100061:      89 04 24         mov     %eax,(%esp)
20 f0100064:      e8 d7 ff ff ff   call    f0100040 <test_backtrace>
21 f0100069:      eb 1c            jmp     f0100087 <test_backtrace+0x47>
22      else
23      mon_backtrace(0, 0, 0);
24 f010006b:      c7 44 24 08 00 00 00 movl    $0x0,0x8(%esp)
25 f0100072:      00
26 f0100073:      c7 44 24 04 00 00 00 movl    $0x0,0x4(%esp)
27 f010007a:      00
28 f010007b:      c7 04 24 00 00 00 00 movl    $0x0,(%esp)
29 f0100082:      e8 18 07 00 00   call    f010079f <mon_backtrace>
30      cprintf("leaving test_backtrace %d\n", x);
31 f0100087:      89 5c 24 04      mov     %ebx,0x4(%esp)
32 f010008b:      c7 04 24 fc 18 10 f0 movl    $0xf01018fc,(%esp)
33 f0100092:      e8 9a 08 00 00   call    f0100931 <cstdio>
34 }

```

<https://github.com/csnlp/MIT6828>

5 Reference

1. bysui's github and blog

- github: <https://github.com/bysui/mit6.828>
- blog: <https://blog.csdn.net/bysui>

2. SmallPond's github and blog

- github: https://github.com/SmallPond/MIT6.828_OS
- blog: https://me.csdn.net/Small_Pond

3. SimpCosm's github

- github: <https://github.com/SimpCosm/6.828>

4. fatsheep9146's blog

- blog: <https://www.cnblogs.com/fatsheep9146/category/769143.html>

<https://github.com/csnlp/MIT6828>